



Core C++ 2024

Open Is Good

YOMM2: Fast, Orthogonal Open (Multi) Methods

Jean-Louis Leroy - jl@leroy.nyc

**IF YOMM2 IS THE
SOLUTION,
WHAT IS THE
PROBLEM?**

THE EXPRESSION PROBLEM

in a polymorphic system...

- existing operations += new types?
- existing types += new operations?

should be possible, easy

C++ COMPILE-TIME POLYMORPHISM

(aka templates)

- existing operations += new types? **easy**
- existing types += new operations? **easy**

to wit: the STL

C++ RUN-TIME POLYMORPHISM

- existing operations += new types?
easy: virtual functions, derivation
- existing types += new operations?
emmmm...

CASE STUDY

ABSTRACT SYNTAX TREE

```
struct Node {
    virtual ~Node() {}
    virtual int value() const = 0;
};

struct Number : Node {
    explicit Number(int value) : val(value) { }
    int value() const override { return val; }
    int val;
};

struct Plus : Node {
    Plus(const Node& left, const Node& right) : left(left), right(right) { }
    int value() const override { return left.value() + right.value(); }
    const Node& left; const Node& right;
};

struct Times : Node {
    Times(const Node& left, const Node& right) : left(left), right(right) { }
    int value() const override { return left.value() * right.value(); }
    const Node& left; const Node& right;
};
```

AST

```
int main() {  
    Number n2(2), n3(3), n4(4);  
    Plus sum(n3, n4);  
    Times product(n2, sum);  
  
    const Node& expr = product;  
    cout << expr.value() << "\n";  
  
    return 0;  
}
```

Output:

ADD AN OPERATION

```
cout << to_rpn(expr) << " = " << expr->value() << "\n";  
//      ^^^^^^
```

Output:

```
2 3 4 * + = 14
```

VIRTUAL FUNCTION?

```
struct Node {  
    // as before  
    virtual string to_rpn() const = 0;  
};  
  
struct Number : Node {  
    // as before  
    string to_rpn() const override { return to_string(val); }  
};  
  
struct Plus : Node {  
    // as before  
    string to_rpn() const override { return left.to_rpn() + " " + right.to_rpn() + " +"; }  
};  
  
struct Times : Node {  
    // as before  
    string to_rpn() const override { return left.to_rpn() + " " + right.to_rpn() + " &"; }  
};
```

banana -> gorilla -> jungle
(C) Erlang creator Joe Armstrong

TYPE SWITCH?

```
string to_rpn(const Node& node) {  
    if (auto expr = dynamic_cast<const Number*>(&node)) {  
        return to_string(expr->value());  
    } else if (auto expr = dynamic_cast<const Plus*>(&node)) {  
        return to_rpn(expr->left) + " " + to_rpn(expr->right) + " +";  
    } else if (auto expr = dynamic_cast<const Times*>(&node)) {  
        return to_rpn(expr->left) + " " + to_rpn(expr->right) + " *";  
    }  
    throw runtime_error("unknown node type");  
}
```

- operations += types: nope

VISITOR?

```
struct Node {  
    // as before  
    struct Visitor {  
        virtual void accept(const Number& expr) = 0;  
        virtual void accept(const Plus& expr) = 0;  
        virtual void accept(const Times& expr) = 0;  
    };  
    virtual void visit(Visitor& viz) const = 0;  
};  
  
struct Number : Node {  
    // as before  
    void visit(Visitor& viz) override { viz.accept(*this); }  
};  
  
struct Plus : Node {  
    void visit(Visitor& viz) override { viz.accept(*this); }  
};  
// etc.
```

VISITOR...

```
struct RPNVisitor : Node::Visitor {
    void accept(const Number& expr) {
        result = to_string(expr.val);
    }
    void accept(const Plus& expr) {
        expr.left.visit(*this);
        string l = result;
        expr.right.visit(*this);
        result = l + " " + result + " +";
    }
    void accept(const Times& expr) { ... }
    string result;
};

string to_rpn(const Node& node) {
    RPNVisitor viz;
    node.visit(viz);
    return viz.result;
}
```

- a lot of work
- more visitors, or more complexity (non-const...)
- operations += types: nope

FUNCTION TABLE?

```
unordered_map<type_index, string (*)(const Node&)> RPNformatters;

string to_rpn(const Node& node) {
    return RPNformatters[typeid(node)](node);
}

struct Init {
    Init() {
        RPNformatters[typeid(Number)] = [](const Node& node) {
            return to_string(static_cast<const Number&>(node).val); };
        RPNformatters[typeid(Plus)] = [](const Node& node) {
            auto expr = static_cast<const Plus&>(node);
            return to_rpn(expr.left) + " " + to_rpn(expr.right) + " +"; };
        RPNformatters[typeid(Times)] = [](const Node& node) {
            auto expr = static_cast<const Times&>(node);
            return to_rpn(expr.left) + " " + to_rpn(expr.right) + " *"; };
    }
} init;
```

- types += operations: yes
- operations += types: yes

POLL

Only one vote!

1. virtual function
2. type switch
3. visitor
4. function table
5. they all stink

OPEN METHODS

OPEN METHODS

- free virtual functions
 - i.e. virtual functions that exist outside of a class
- existing types += new operations

YOMM2 OPEN METHODS

```
struct Node {  
    virtual string to_rpn(/* const Node* */) const = 0;  
};
```

```
declare_method(string, to_rpn, (virtual_<const Node>));
```

Common Lisp: defgeneric, Clojure: defmulti

```
struct Plus : Node {  
    string to_rpn(/* const Node* this */) const override {  
        return left.to_rpn() + " " + right.to_rpn() + " +";  
    }  
};
```

```
define_method(string, to_rpn, (const Plus& expr)) {  
    return to_rpn(expr.left) + " " + to_rpn(expr.right) + " +";  
}
```

Common Lisp, Clojure: defmethod

AST

```
#include <yorel/yomm2/keywords.hpp>

declare_method(string, to_rpn, (virtual_<const Node&>));

define_method(string, to_rpn, (const Number& expr)) {
    return std::to_string(expr.val);
}

define_method(string, to_rpn, (const Plus& expr)) {
    return to_rpn(expr.left) + " " + to_rpn(expr.right) + " +";
}

define_method(string, to_rpn, (const Times& expr)) {
    return to_rpn(expr.left) + " " + to_rpn(expr.right) + " *";
}

register_classes(Node, Number, Plus, Times);

int main() {
    yorel::yomm2::update();
    cout << to_rpn(expr) << " = " << expr.value() << "\n";
    return 0;
}
```

AST: WHAT ABOUT VALUE?

- value in the node hierarchy => interpreter
- AST classes should *only* represent the tree

```
declare_method(int, value, (virtual_<Node&>));

define_method(int, value, (Number& expr)) {
    return expr.val;
}

define_method(int, value, (Plus& expr)) {
    return value(expr.left) + value(expr.right);
}
```

PERFORMANCE

```
mov    rax, qword ptr [rdi]
mov    rdx, qword ptr [rip+fast_perfect_hash<release>::mult]

imul   rdx, qword ptr [rax-8]
mov    cl, byte ptr [rip+fast_perfect_hash<release>::shift]

shr    rdx, cl
mov    rax, qword ptr [rip+vp_ptr_vector<release>::vp_ptr]

mov    rax, qword ptr [rax+8*rdx]
mov    rcx, qword ptr [rip+method<value, int (virtual_<Node const&>)::fn+80]

jmp    qword ptr [rax+8*rcx]
```

- llvm-mca: 4.2 cycles
- benchmarks (google, rdtsc): 15-65% slower than equivalent native virtual function call (using perfect integer hash; but see `virtual_ptr`)
- Optimizing Away C++ Virtual Functions May Be Pointless - Shachar Shemesh - CppCon 2023

MULTIPLE DISPATCH

sometimes useful

```
add(Matrix, Matrix)           -> Matrix
                                add all elements
add(DiagonalMatrix, DiagonalMatrix) -> DiagonalMatrix
                                just add diagonals

fight(Human, Creature, Axe)    -> not agile enough to wield
fight(Warrior, Creature, Axe)  -> chop it into pieces
fight(Warrior, Dragon, Axe)    -> die a honorable death
fight(Human, Dragon, Hands)    -> congratulations! you have just
                                vanquished a dragon with your
                                bare hands! (unbelievable,
                                isn't it?)
```

- works just like selecting from set of overloads (but at runtime!)
- ambiguities can arise

POLL

IS THIS OOP?

Only one vote!

1. Yes
2. No

POLL

WHAT DO YOU PREFER?

Only one vote!

1. virtual function, type switch, visitor, function table
2. open methods

INSIDE YOMM2

INSIDE YOMM2

- pure C++17 (no extra tooling)
- "constant" time dispatch (proportional to #vargs)
- uses tables of function pointers
- object -> dispatch data?
 - perfect integer hash of &type_info

A PAYROLL APPLICATION

- *Role*
 - Employee
 - Manager
 - Founder
- *Expense*
 - Cab, Jet
 - *Public*
 - Bus, Train

THE pay UNI- METHOD

```
declare_method(double, pay, (virtual_<Employee&>));  
  
define_method(double, pay, (Employee&)) {  
    return 3000;  
}  
  
define_method(double, pay, (Manager& manager)) {  
    return next(manager) + 2000;  
}
```

DECLARE_METHOD

```
declare_method(double, pay, (virtual_<Employee&>));
```

```
struct YoMm2_S_pay;
```

```
inline double  
pay(yomm2::detail::remove_virtual<virtual_<const Employee&>> a0) {  
    return yomm2::method<  
        YoMm2_S_pay, double(virtual_<const Employee&>),  
        yomm2::default_policy>::  
        fn(std::forward<  
            yomm2::detail::remove_virtual<virtual_<const Employee&>>>(a0));  
};
```

```
yomm2::method<  
    YoMm2_S_pay, double(virtual_<const Employee&>),  
    yomm2::default_policy>  
pay_yOMM2_selector_  
    yomm2::detail::remove_virtual<virtual_<const Employee&>> a0);
```

DEFINE_METHOD

```
define_method(double, pay, (Employee&)) { return 3000; }
```

```
namespace { namespace YoMm2_gS_10 {  
template<typename T> struct _yOMM2_select;  
template<typename... A> struct _yOMM2_select<void(A...)> {  
    using type = decltype(pay_yOMM2_selector_(std::declval<A>()...));  
};  
using _yOMM2_method = _yOMM2_select<void(const Employee&)>::type;  
using _yOMM2_return_t = _yOMM2_method::return_type;  
_yOMM2_method::function_pointer_type next;  
struct _yOMM2_spec {  
    static YoMm2_gS_10::_yOMM2_method::return_type  
    yOMM2_body(const Employee&);  
};  
_yOMM2_method::add_function<_yOMM2_spec::yOMM2_body>  
    YoMm2_gS_11(&next, typeid(_yOMM2_spec).name()); } }  
YoMm2_gS_10::_yOMM2_method::return_type  
YoMm2_gS_10::_yOMM2_spec::yOMM2_body(const Employee&) {  
    return 3000;  
}
```

UPDATE

- uses class and method info registered by static ctors
- builds a representation of class hierarchies
- builds dispatch tables
- finds a perfect (not minimal) hash function for the `type_infos`
 - $H(x) = (M * x) \gg S$

DISPATCHING A UNI-METHOD

- pretty much like virtual member functions
- method table contains a pointer to the effective function
- only it is not at a fixed offset in the method table

DISPATCHING A UNI-METHOD

during update

```
method<pay>::slots_strides[] = { 1 };  
  
// method table for Employee  
mtbls[ H(&typeid(Employee)) ] = {  
    ..., // used by approve,  
    wrapper(pay(Employee&))  
};  
  
// method table for Manager  
mtbls[ H(&typeid(Manager&)) ] = {  
    ..., // used by approve,  
    wrapper(pay(Manager&))  
};
```

DISPATCHING A UNI-METHOD

```
pay(bill)
```

=>

```
mtbls[ H(&typeid(bill)) ] // mtable for type  
[ method<pay>::slots_strides[0] ] // pointer to fun  
(bill) // call
```

ASSEMBLER

```
double call_pay(Employee& e) { return pay(e); }
```

```
mov rax, qword ptr [rdi]           ; vptr
mov rdx, qword ptr [rip + hash_mult] ; M
imul rdx, qword ptr [rax - 8]       ; M * &typeid(e)
movzx ecx, byte ptr [rip + hash_shift] ; S
shr rdx, cl                         ; >> S
mov rax, qword ptr [rip + vptrs]    ; vptrs
mov rax, qword ptr [rax + 8*rdx]     ; vptr
mov rcx, qword ptr [rip + slots_strides] ; slot
jmp qword ptr [rax + 8*rcx]
```

approve **MULTI-METHOD**

```
declare_method(bool, approve,  
    (virtual_<Role&>, virtual_<Expense&>, double));  
  
define_method(bool, approve, (Role& r, Expense& e, double amount)) {  
    return false;  
}  
  
define_method(bool, approve, (Employee& r, Public& e, double amount)) {  
    return true;  
}  
  
define_method(bool, approve, (Manager& r, Taxi& e, double amount)) {  
    return true;  
}  
  
define_method(bool, approve, (Founder& r, Expense& e, double amount)) {  
    return true;  
}
```

DISPATCHING A MULTI-METHOD

- it's a little more complicated
- uses a multi-dimensional dispatch table
- size can grow very quickly
- table must be "compressed", devoid of redundancies
- in fact the "uncompressed" table never exists
- works in terms of class *groups*, not classes

DISPATCHING A MULTI-METHOD

	Expense+Jet	Public+Bus+Train	Cab
Role	R,X	R,X	R,X
Employee	R,X	E,P	R,X
Manager	R,X	E,P	M,C
Founder	F,X	F,X	F,X

(column major)

BUILDING THE DISPATCH TABLE

- Fast Algorithms for Compressed Multi-Method Dispatch, Eric Amiel, Eric Dujardin, Eric Simon, 1996
- Open Multi-Methods for C++11, Part 3 - Inside Yomm11: Data Structures and Algorithms, Jean-Louis Leroy, 2013

DISPATCHING A MULTI-METHOD

```
method<approve>::slots_strides = { 0, 4, 0 };

mtbls[ H(&typeid(Employee)) ] = {
    // & of (Employee,Expense+Jet) cell
    // used by pay
};

mtbls[ H(&typeid(Manager)) ] = {
    // & of (Manager,Expense+Jet) cell
    // used by pay
};

mtbls[ H(&typeid(Expense)) ] = { 0 }; // also for Jet
mtbls[ H(&typeid(Public)) ] = { 1 }; // also for Bus, Train
mtbls[ H(&typeid(Cab)) ] = { 2 };
```


DISPATCHING A MULTI-METHOD

```
approve(bill, ticket, 6)
```

=>

```
std::uintptr_t* slots_strides = method<approve>::slots_strides;

mtbls[ H(&typeid(bill)) ]           // method table for bill
[ slots_strides[0] ]                // ptr to cell in 1st column
[ mtbls [ H(&typeid(ticket)) ] ]    // method table for ticket
[ slots_strides[2] ]                // column
* slots_strides[1]                  // stride
]                                    // pointer to function
(bill, ticket, 6)                    // call
```

YOMM2 VS OTHER SYSTEMS

- Pirkelbauer - Solodkyi - Stroustrup (PSS)
- Cmm
- Loki / Modern C++

YOMM2 VS PSS

- Solodkyi's papers on open methods etc.:
 - Open Multi-Methods for C++
 - Design and Evaluation of C++ Open Multi-Methods
 - Simplifying the Analysis of C++ Programs
- PSS attempts harder to resolve ambiguities
- YOMM2 overrides not visible as overloads, cannot specialize multiple methods
- YOMM2 supports smart pointers, next

EVOLUTION OF YOMM2

PAST

PAST

PAST

- goals:
 - help promote adoption in the language
 - submit to Boost
 - talk about it (CppCon 2018...)
 - 2018-2020: only bug fixes, cleanup...
- results:
 - Boost community: no interest
 - standard committee: no interest

PRESENT

- 2020: give up on adoption in the standard
- new features
- `virtual_ptr`

VIRTUAL_PTR

```
declare_method(int, value, (virtual_ptr<Node>));
```

```
int call_via_vptr(virtual_ptr<const Node> node) {  
    return value(node);  
}
```

```
mov rax, qword ptr [rip + method<value, int (virtual_ptr<Node>)::fn+80]  
mov rax, qword ptr [rsi + 8*rax]  
jmp rax
```


VIRTUAL_PTR

```
auto make_node_ptr(const Node& node) {  
    return virtual_ptr(node);  
}
```

```
mov rax, rdi  
mov rcx, qword ptr [rdi]  
mov rdx, qword ptr [rcx - 8]  
lea rcx, [rip + typeinfo for Node]  
cmp rdx, rcx  
je .LBB7_1  
imul rdx, qword ptr [rip + fast_perfect_hash<release>::hash_mult]  
movzx ecx, byte ptr [rip + fast_perfect_hash<release>::hash_shift]  
shr rdx, cl  
shl rdx, 3  
add rdx, qword ptr [rip + vptr_vector<release>::vptrs]  
mov rdx, qword ptr [rdx]  
ret  
.LBB7_1:  
lea rdx, [rip + method_tables<release>::static_vptr<Node>]  
mov rdx, qword ptr [rdx]  
ret
```

VIRTUAL_PTR

```
auto make_final_node_ptr(const Plus& node) {  
    return final_virtual_ptr(node);  
}
```

classes need not be polymorphic

```
mov rax, rdi  
mov rdx, qword ptr [rip + method_tables<release>::static_vptr<Plus>]  
ret
```

NEW FEATURES

- virtual_ptr
- core API

CORE API

```
struct value_id;
using value = method<value_id, int(virtual_<const Node&>>>;

auto result = value::fn(expr);

int number_value(const Number& node) {
    return node.val;
}
value::add_function<number_value> add_number_value;

template<class NodeClass, class Op>
struct binary_value {
    static int fn(const NodeClass& expr) {
        return Op()(value::fn(expr.left), value::fn(expr.right));
    }
};

YOMM2_STATIC(value::add_definition<binary_value<Plus, std::plus<int>>>>);
YOMM2_STATIC(value::add_definition<binary_value<Times, std::multiplies<int>>>>);

YOMM2_STATIC(use_classes<Node, Number, Plus, Times>);
```

NEW FEATURES

- virtual_ptr
- core API
- header only
- friendship
- member methods
- policies and facets
 - custom RTTI
 - custom error handling, trace, vptr placement...
- match (beat?) virtual function speed
- pre-calculate dispatch tables
- vcpkg, Conan, Compiler Explorer

FUTURE

- Boost.OpenMethod ?
- C++20
- C++26 (reflection)

Q&A

GitHub:



examples are on Compiler Explorer:

<https://jll63.github.io/yomm2/ce/slides.html>

(redirects to volatile godbolt.org short URL)